Confido Health — Product Manager, Platform Take-Home

Two parts: in **Part 1** you build a voice agent for a clinic; in **Part 2** you use what you learned to make the case for how Confido would productize agent-building.

Part 1 — Build the agent

Just like a Confido FDE standing up a new clinic, you get two things (attached): the clinic's **rules** and some **sample data**. Build a voice agent for the clinic that handles inbound calls for:

- **Appointment management** — book (new and existing patients), reschedule, cancel, confirm
- **Clinic FAQs** — hours, locations, parking
- **Transfers** — to the front desk, and to a medical professional for urgent medical attention

Handle the obvious failure cases gracefully (no matching patient or appointment, an unrecognized request, a caller who wants a human).

Minimum: your agent makes tool calls that return mock data from the sample set. Getting those tools to fire live via a small hosted stub + tunnel (e.g. `ngrok`) is a nice-to-have if you have time — not required. **Retell preferred** (Vapi, Bland, LiveKit, Pipecat are fine). Spend your time on the rules, the design, and iteration — not on telephony.

Submit:

1. **The agent** — a Retell workspace / agent we can open and view.
2. **Architecture write-up** — your agent's structure, guardrails, and the design principles behind your choices and why you made them. For example: how you decided what belongs in the prompt vs. a tool call vs. a knowledge base. We're looking for that kind of principled decision-making across the build.
3. **A high-level iteration log** — what you changed and why as you built.

Part 2 — Productizing agent-building

A short document: based on what you just did, how would you approach turning this one-off into a platform way of **building, refining, and iterating** these agents across many clinics — faster and at a higher quality bar?

Briefly cover:

- **Where an agentic (AI-driven) solve fits** in the build / iteration process — and where it doesn't.
- **Experiments** — what you'd run, the metric each would move, and how you'd sequence them. Prototypes tied to your experiments are highly preferred.
- **Prioritization** — what you'd do first, and why.

Optional but highly preferred — a short video

5–10 minutes walking us through your Part 1 and Part 2 answers and why you made the calls you did.